Deep contextualized word representation

AntNLP

Yupei Du

Language Models

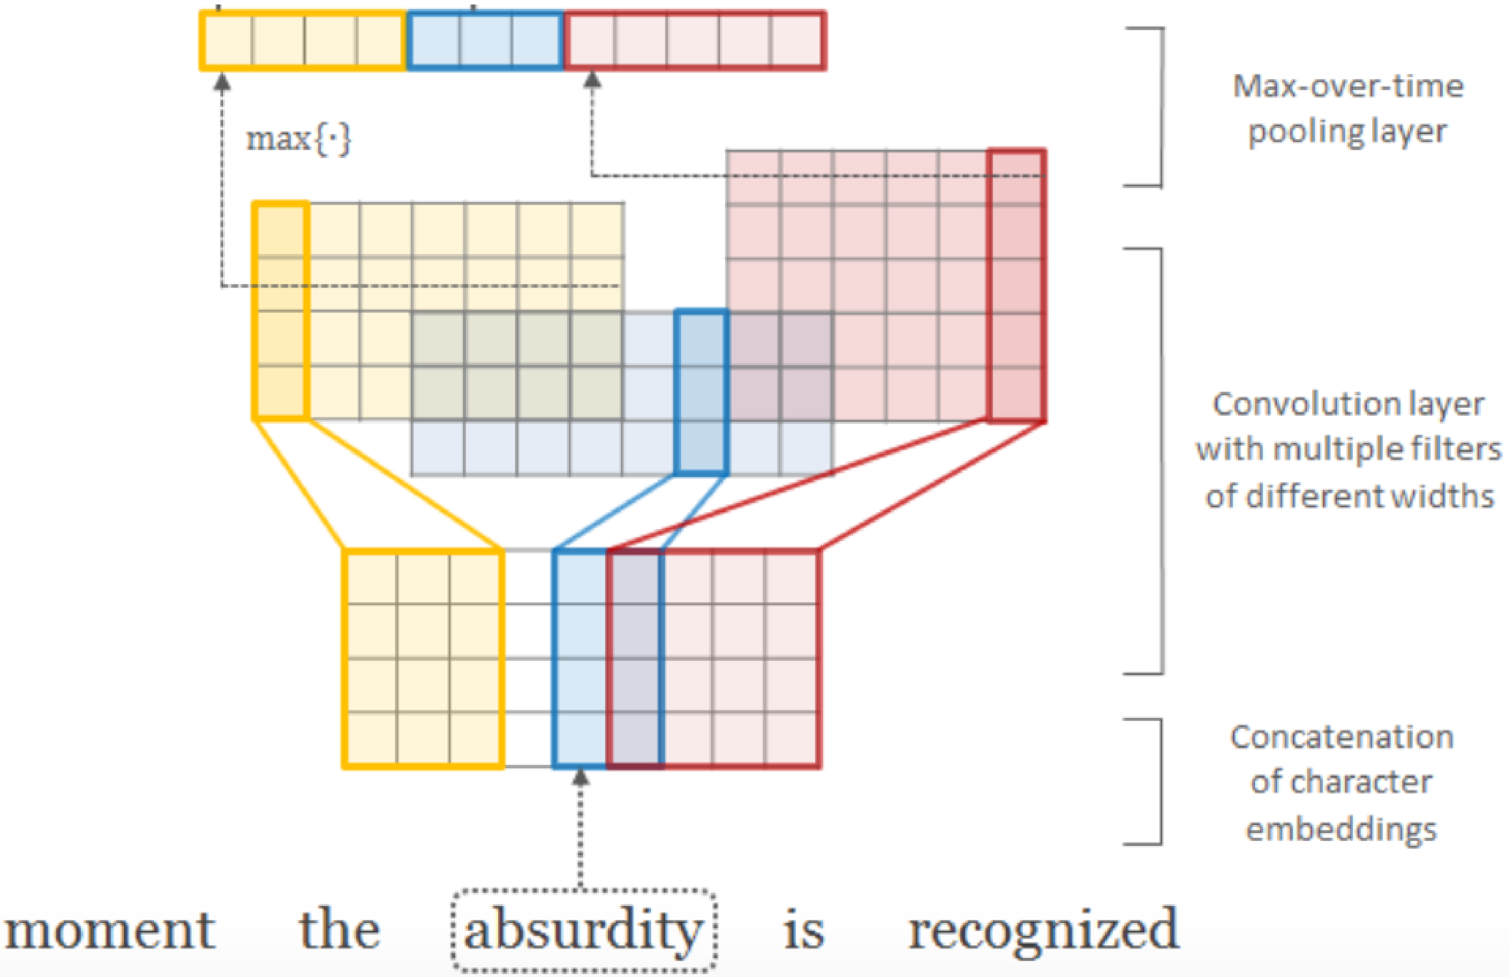
max{·}
Max-over-time pooling layer
Convolution layer with multiple filters of different widths
Concatenation of character embeddings
moment
the
absurdity
is
recognized

Language Models

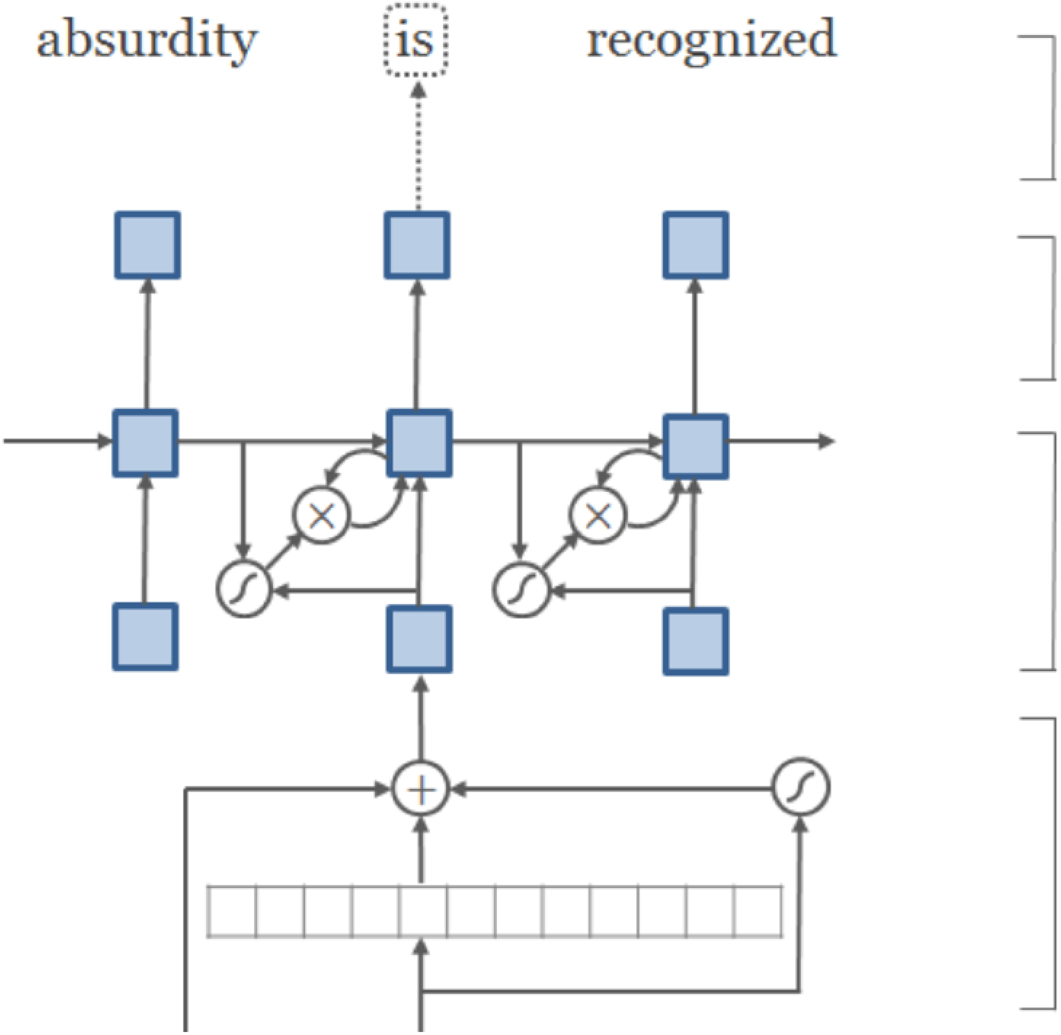

Cross entropy loss between next word and prediction

Softmax output to obtain distribution over next word

Long short-term memory network

Highway network

ELMo

Deep-BiLSTM

$$\begin{aligned}\sum_{k=1}^{N} (\log p(t_k \mid t_1, \ldots, t_{k-1}; \Theta_x, \overrightarrow{\Theta}_{LSTM}, \Theta_s) \\ + \log p(t_k \mid t_{k+1}, \ldots, t_N; \Theta_x, \overleftarrow{\Theta}_{LSTM}, \Theta_s) \,)\end{aligned}$$

ELMo

$$\begin{aligned} R_k &= \{\mathbf{x}_k^{LM}, \overrightarrow{\mathbf{h}}_{k,j}^{LM}, \overleftarrow{\mathbf{h}}_{k,j}^{LM} \mid j = 1, \ldots, L\} \\ &= \{\mathbf{h}_{k,j}^{LM} \mid j = 0, \ldots, L\}, \end{aligned}$$

$$\mathbf{ELMo}_k^{task} = E(R_k; \Theta^{task}) = \gamma^{task} \sum_{j=0}^{L} s_j^{task} \mathbf{h}_{k,j}^{LM}.$$

Experiments

SQuAD(Question answering) : 100K+ crowd sourced question-answer pairs where the answer is a span in a given Wikipedia paragraph

SNLI(Textual entailment) : 550K hypothesis/premise pairs.determining whether a "hypothesis" is true, given a "premise"

SRL(Semantic role labeling) : models the predicate-argument structure of a sentence, and is often described as answering "Who did what to whom"

Coref(Coreference resolution) : clustering mentions in text that re- fer to the same underlying real world entities

Experiments

NER(Named entity extraction) : tag entities with four different entity types

SST-5(Sentiment analysis) : selecting one of five labels (from very negative to very positive) to describe a sentence from a movie review

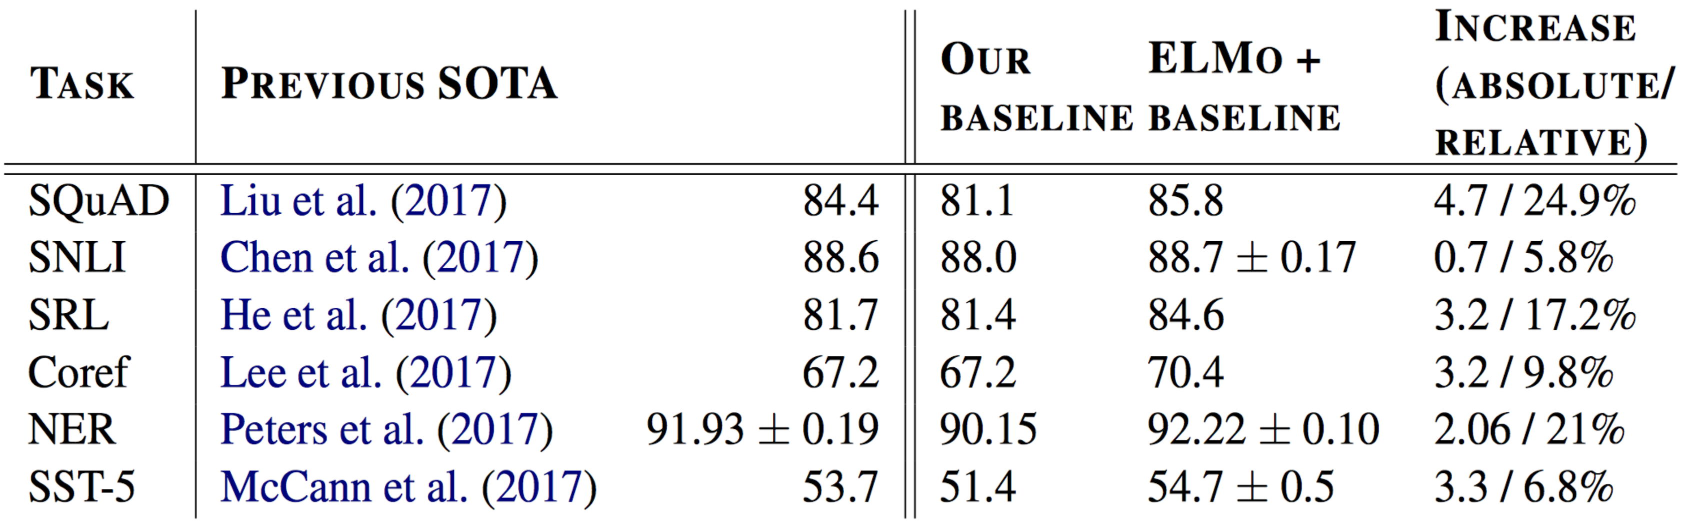

TASK	PREVIOUS SOTA		OUR BASELINE	ELMO + BASELINE	INCREASE (ABSOLUTE/ RELATIVE)
SQuAD	Liu et al. (2017)	84.4	81.1	85.8	4.7 / 24.9%
SNLI	Chen et al. (2017)	88.6	88.0	88.7 ± 0.17	0.7 / 5.8%
SRL	He et al. (2017)	81.7	81.4	84.6	3.2 / 17.2%
Coref	Lee et al. (2017)	67.2	67.2	70.4	3.2 / 9.8%
NER	Peters et al. (2017)	91.93 ± 0.19	90.15	92.22 ± 0.10	2.06 / 21%
SST-5	McCann et al. (2017)	53.7	51.4	54.7 ± 0.5	3.3 / 6.8%

Analysis

Regularize

Adding $\lambda||w||_2^2$ to loss function

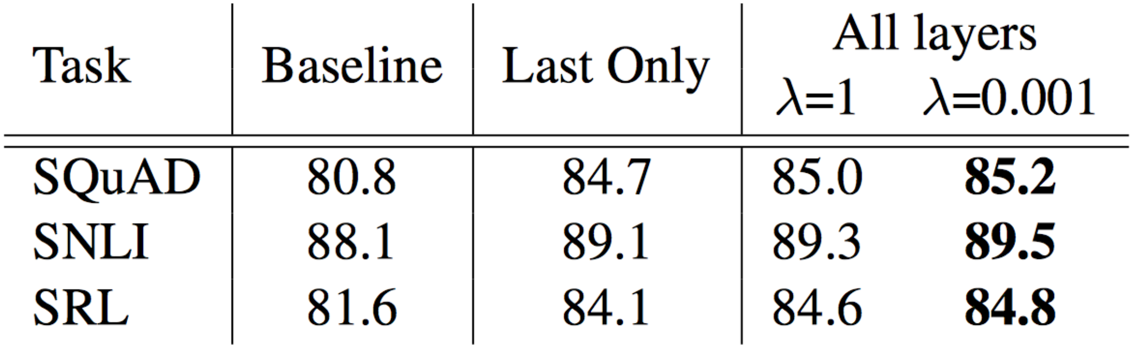

Task	Baseline	Last Only	All layers	
			λ=1	λ=0.001
SQuAD	80.8	84.7	85.0	**85.2**
SNLI	88.1	89.1	89.3	**89.5**
SRL	81.6	84.1	84.6	**84.8**

Input or Output?

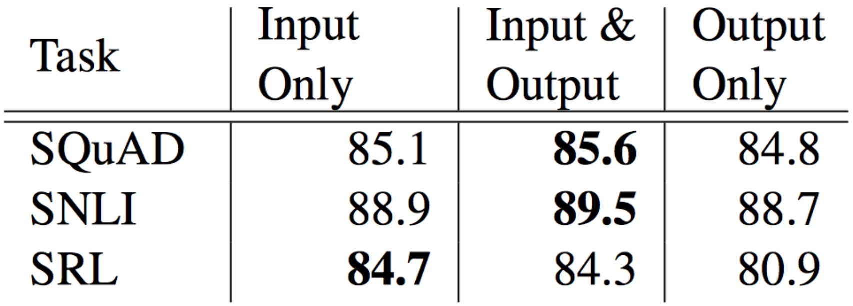

Task	Input Only	Input & Output	Output Only
SQuAD	85.1	**85.6**	84.8
SNLI	88.9	**89.5**	88.7
SRL	**84.7**	84.3	80.9

Analysis

What information is captured?

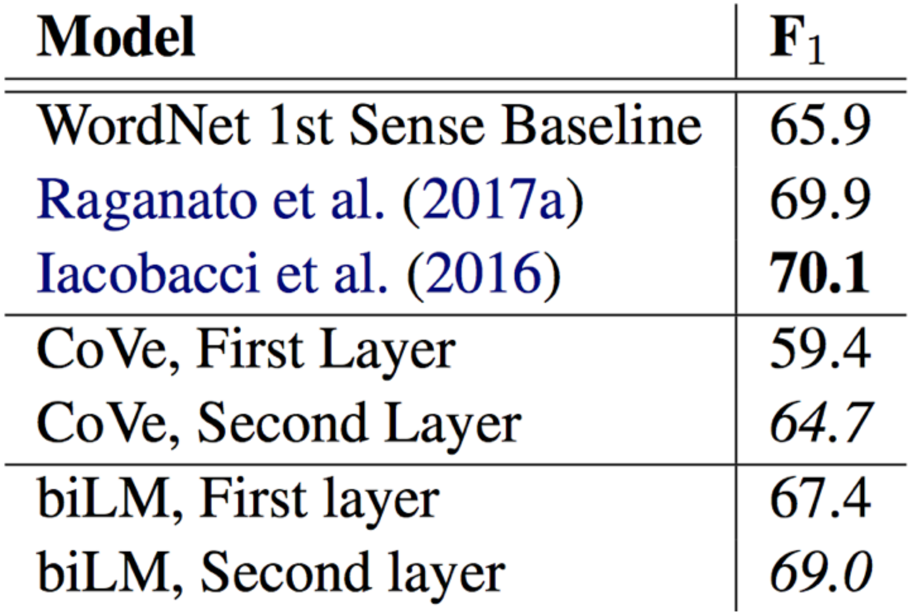

Model	F_1
WordNet 1st Sense Baseline	65.9
Raganato et al. (2017a)	69.9
Iacobacci et al. (2016)	**70.1**
CoVe, First Layer	59.4
CoVe, Second Layer	*64.7*
biLM, First layer	67.4
biLM, Second layer	*69.0*

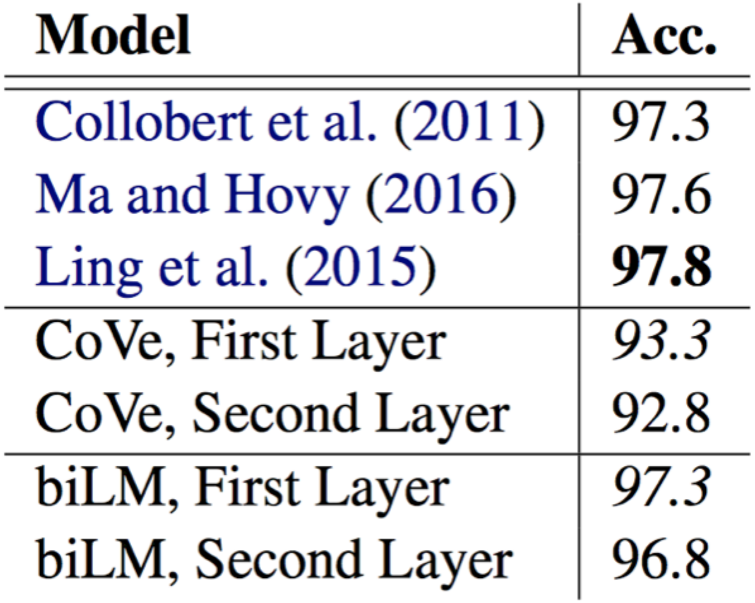

Model	Acc.
Collobert et al. (2011)	97.3
Ma and Hovy (2016)	97.6
Ling et al. (2015)	**97.8**
CoVe, First Layer	*93.3*
CoVe, Second Layer	92.8
biLM, First Layer	*97.3*
biLM, Second Layer	96.8

Analysis

Sample efficiency

Speed

$$486 \rightarrow 10\ epochs\ in\ SRL$$

dataset size

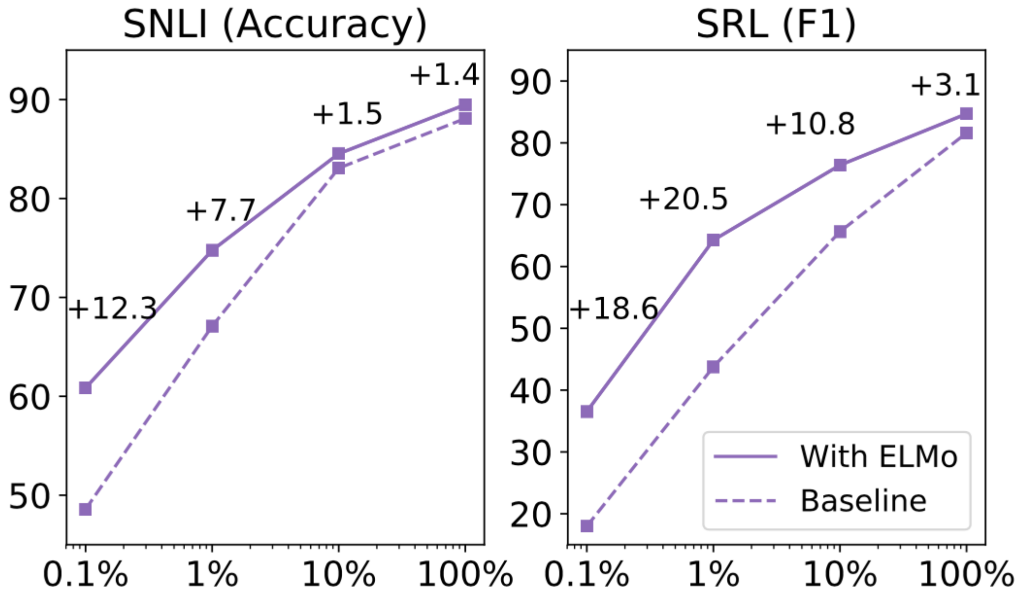

Analysis

Visualization of weights

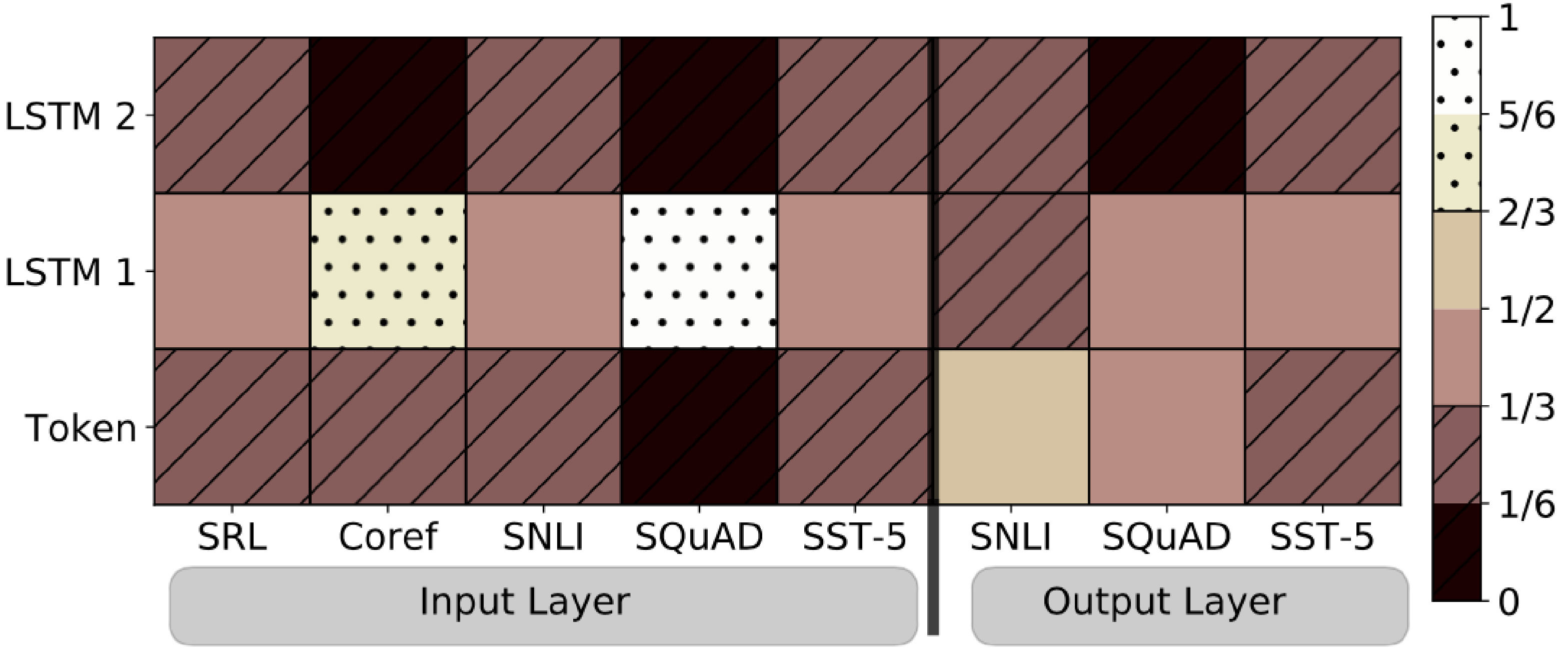

Thank you!

Q&A